

## ASSIGNMENT 1 MACHINE LEARNING

SYABAB ALHAQQI BIN JAAFAR 2211117

### Q1 housing

A) Parameters selected is: ['longitude', 'latitude', 'housing\_median\_age', 'total\_rooms', 'total\_bedrooms', 'population', 'households', 'median\_income']. These are all the important parameters to ensure accurate data prediction.

#### Coding:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, r2_score

# Load Dataset
df = pd.read_csv("housing.csv")

# Data Preparation
# Handle missing values
df['total_bedrooms'].fillna(df['total_bedrooms'].median(), inplace=True)

# One-Hot Encoding for categorical variable
df = pd.get_dummies(df, columns=['ocean_proximity'], drop_first=True)

# Feature Scaling
scaler = StandardScaler()

numerical_features = ['longitude', 'latitude', 'housing_median_age', 'total_rooms',
                      'total_bedrooms', 'population', 'households', 'median_income']
```

```
df[numerical_features] = scaler.fit_transform(df[numerical_features])
```

```
# Define Features and Target
```

```
X = df.drop(columns=['median_house_value'])
```

```
y = df['median_house_value']
```

```
# Train-test split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# --- Multiple Linear Regression ---
```

```
lr_model = LinearRegression()
```

```
lr_model.fit(X_train, y_train)
```

```
y_lr_pred = lr_model.predict(X_test)
```

```
lr_mse = mean_squared_error(y_test, y_lr_pred)
```

```
lr_r2 = r2_score(y_test, y_lr_pred)
```

```
# --- Polynomial Regression (Degree 2) ---
```

```
poly = PolynomialFeatures(degree=2)
```

```
X_train_poly = poly.fit_transform(X_train)
```

```
X_test_poly = poly.transform(X_test)
```

```
poly_model = LinearRegression()
```

```
poly_model.fit(X_train_poly, y_train)
```

```
y_poly_pred = poly_model.predict(X_test_poly)
```

```
poly_mse = mean_squared_error(y_test, y_poly_pred)
```

```
poly_r2 = r2_score(y_test, y_poly_pred)
```

```
# --- Support Vector Regression (SVR) ---
```

```
scaler_y = StandardScaler()
```

```
y_train_scaled = scaler_y.fit_transform(y_train.values.reshape(-1, 1)).ravel()
```

```
y_test_scaled = scaler_y.transform(y_test.values.reshape(-1, 1)).ravel()
```

```

svr_model = SVR(kernel='rbf', C=100, gamma=0.1, epsilon=0.1)
svr_model.fit(X_train, y_train_scaled)
y_svr_pred_scaled = svr_model.predict(X_test)
y_svr_pred = scaler_y.inverse_transform(y_svr_pred_scaled.reshape(-1, 1)).ravel()
svr_mse = mean_squared_error(y_test, y_svr_pred)
svr_r2 = r2_score(y_test, y_svr_pred)

# Display results
results = pd.DataFrame({
    "Model": ["Linear Regression", "Polynomial Regression (Degree 2)", "Support Vector
Regression"],
    "MSE": [lr_mse, poly_mse, svr_mse],
    "R2 Score": [lr_r2, poly_r2, svr_r2]
})

print(results)

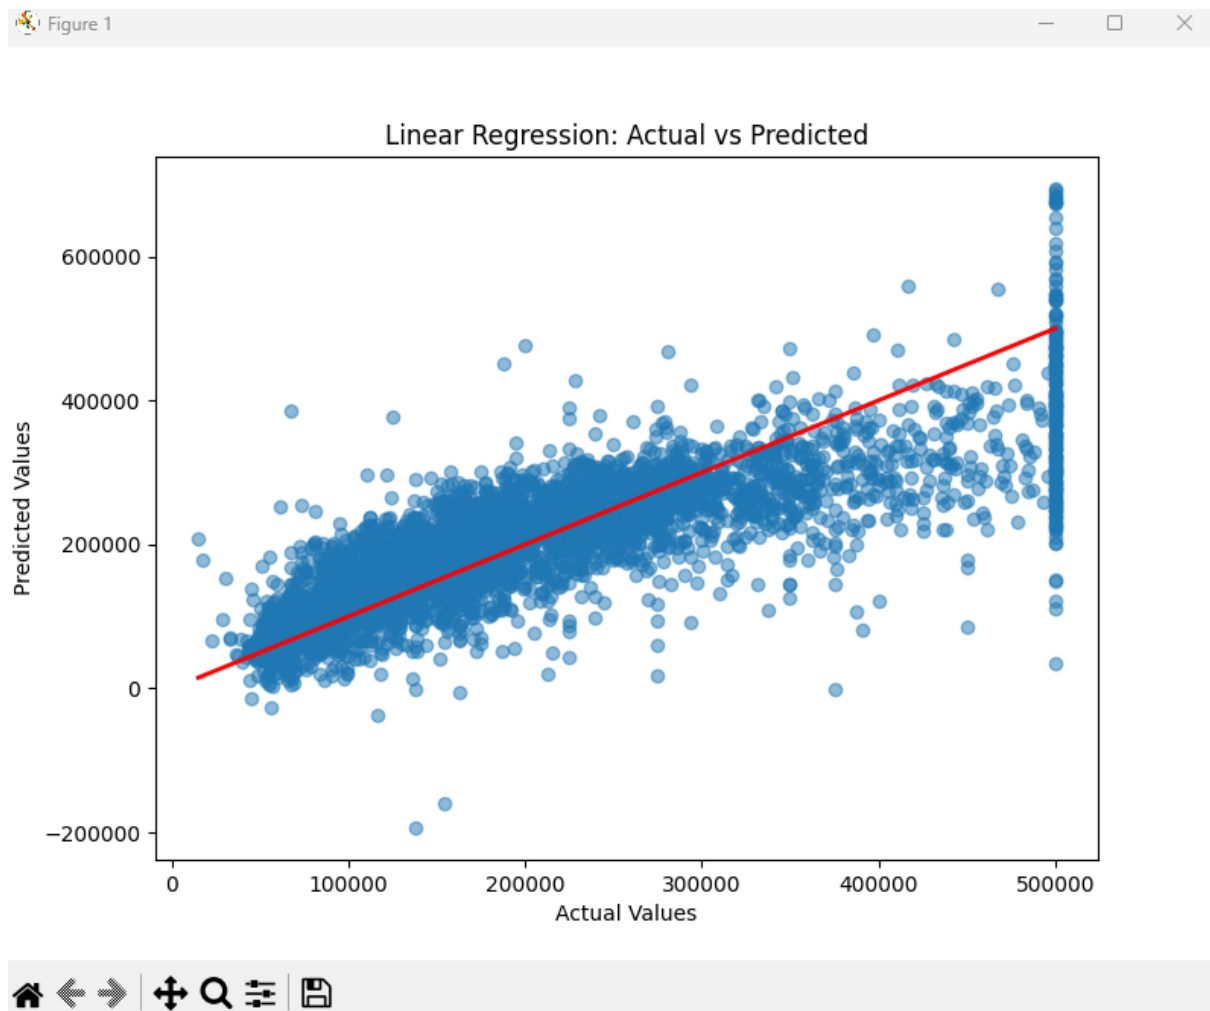
# --- Plot Predictions vs Actual Values ---
def plot_results(y_test, y_pred, title):
    plt.figure(figsize=(8, 6))
    plt.scatter(y_test, y_pred, alpha=0.5)
    plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r', linewidth=2)
    plt.xlabel('Actual Values')
    plt.ylabel('Predicted Values')
    plt.title(title)
    plt.show()

plot_results(y_test, y_lr_pred, "Linear Regression: Actual vs Predicted")
plot_results(y_test, y_poly_pred, "Polynomial Regression: Actual vs Predicted")
plot_results(y_test, y_svr_pred, "SVR: Actual vs Predicted")


```

**F)** The best model is SVR as it has the highest  $R^2$  score and the lowest MAE/MSE.


	Model	MSE	$R^2$ Score
0	Linear Regression	4.908477e+09	0.625424
1	Polynomial Regression (Degree 2)	4.509304e+09	0.655886
2	Support Vector Regression	3.077935e+09	0.765116



Polynomial Regression: Actual vs Predicted



SVR: Actual vs Predicted



## Q2 solar-based energy

**A)** Parameters selected is: ["shortwave\_radiation\_backwards\_sfc", "total\_cloud\_cover\_sfc", "high\_cloud\_cover\_high\_cld\_lay", "medium\_cloud\_cover\_mid\_cld\_lay", "low\_cloud\_cover\_low\_cld\_lay", "temperature\_2\_m\_above\_gnd", "relative\_humidity\_2\_m\_above\_gnd", "mean\_sea\_level\_pressure\_MSL", "total\_precipitation\_sfc", "snowfall\_amount\_sfc", "wind\_speed\_10\_m\_above\_gnd", "wind\_speed\_80\_m\_above\_gnd", "wind\_speed\_900\_mb", "wind\_direction\_10\_m\_above\_gnd", "wind\_direction\_80\_m\_above\_gnd", "wind\_direction\_900\_mb", "angle\_of\_incidence", "zenith", "azimuth"]. These selected features are physically relevant for PV power prediction.

### Coding:

```
import pandas as pd

import matplotlib.pyplot as plt

from sklearn.impute import SimpleImputer

from sklearn.preprocessing import MinMaxScaler, PolynomialFeatures

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.svm import SVR

from sklearn.pipeline import make_pipeline

from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score


# Load the dataset

file_path = "SolarPowerGenerationKaggle_MissingData.xlsx"

df = pd.read_excel(file_path, sheet_name="spg")


# Select relevant features

selected_features = [

    "shortwave_radiation_backwards_sfc", "total_cloud_cover_sfc",

    "high_cloud_cover_high_cld_lay", "medium_cloud_cover_mid_cld_lay",

    "low_cloud_cover_low_cld_lay",

    "temperature_2_m_above_gnd", "relative_humidity_2_m_above_gnd",

    "mean_sea_level_pressure_MSL",
```

```
"total_precipitation_sfc", "snowfall_amount_sfc", "wind_speed_10_m_above_gnd",  
"wind_speed_80_m_above_gnd", "wind_speed_900_mb",  
"wind_direction_10_m_above_gnd",  
"wind_direction_80_m_above_gnd", "wind_direction_900_mb", "angle_of_incidence",  
"zenith", "azimuth"  
]
```

```
df_selected = df[selected_features + ["generated_power_kw"]]
```

```
# Handle missing values with mean imputation
```

```
imputer = SimpleImputer(strategy="mean")
```

```
df_imputed = pd.DataFrame(imputer.fit_transform(df_selected),  
columns=df_selected.columns)
```

```
# Scale data using Min-Max Scaling
```

```
scaler = MinMaxScaler()
```

```
df_scaled = pd.DataFrame(scaler.fit_transform(df_imputed),  
columns=df_selected.columns)
```

```
# Split data into features (X) and target variable (y)
```

```
X = df_scaled.drop(columns=["generated_power_kw"])
```

```
y = df_scaled["generated_power_kw"]
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# --- Multiple Linear Regression ---
```

```
mlr_model = LinearRegression()
```

```
mlr_model.fit(X_train, y_train)
```

```
y_pred_mlr = mlr_model.predict(X_test)
```

```
mae_mlr = mean_absolute_error(y_test, y_pred_mlr)
```

```
mse_mlr = mean_squared_error(y_test, y_pred_mlr)
```

```
r2_mlr = r2_score(y_test, y_pred_mlr)
```

```

# --- Polynomial Regression (Degree 2) ---

poly_degree = 2

poly_model = make_pipeline(PolynomialFeatures(degree=poly_degree),
LinearRegression())

poly_model.fit(X_train, y_train)

y_pred_poly = poly_model.predict(X_test)

mae_poly = mean_absolute_error(y_test, y_pred_poly)

mse_poly = mean_squared_error(y_test, y_pred_poly)

r2_poly = r2_score(y_test, y_pred_poly)


# --- Support Vector Regression (SVR) ---

svr_model = SVR(kernel='rbf')

svr_model.fit(X_train, y_train)

y_pred_svr = svr_model.predict(X_test)

mae_svr = mean_absolute_error(y_test, y_pred_svr)

mse_svr = mean_squared_error(y_test, y_pred_svr)

r2_svr = r2_score(y_test, y_pred_svr)


# Print model evaluation results

print("Multiple Linear Regression:", f"MAE: {mae_mlr:.4f}, MSE: {mse_mlr:.4f}, R2: {r2_mlr:.4f}")

print("Polynomial Regression:", f"MAE: {mae_poly:.4f}, MSE: {mse_poly:.4f}, R2: {r2_poly:.4f}")

print("Support Vector Regression:", f"MAE: {mae_svr:.4f}, MSE: {mse_svr:.4f}, R2: {r2_svr:.4f}")


# --- Plot Predictions vs Actual Values ---

def plot_results(y_test, y_pred, title):

    plt.figure(figsize=(8, 6))

    plt.scatter(y_test, y_pred, alpha=0.5)

    plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r', linewidth=2)

    plt.xlabel('Actual Values')

```



```
plt.ylabel('Predicted Values')
```

```
plt.title(title)
```

```
plt.show()
```

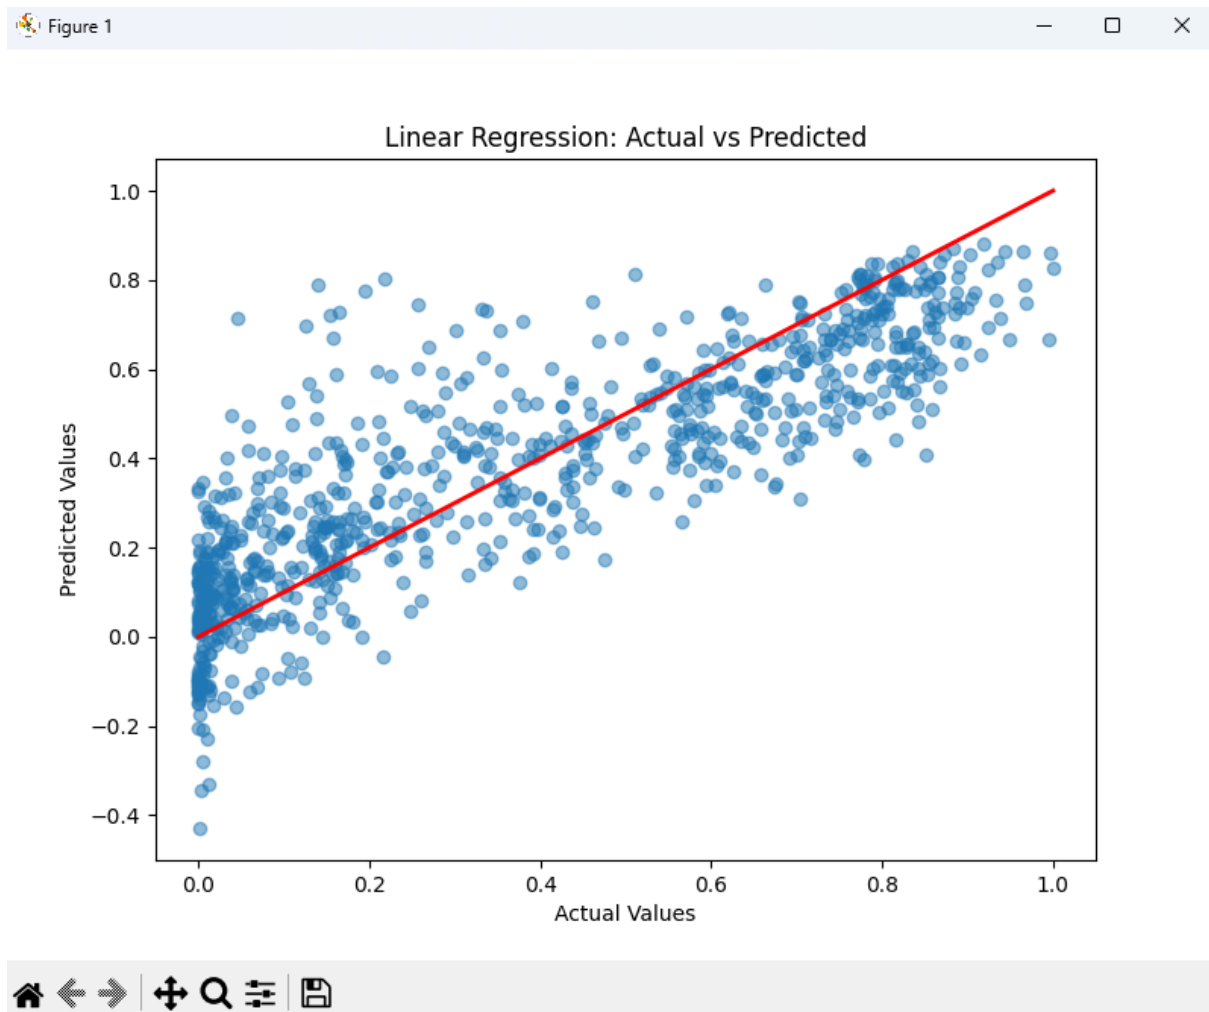
```
plot_results(y_test, y_pred_mlr, "Linear Regression: Actual vs Predicted")
```

```
plot_results(y_test, y_pred_poly, "Polynomial Regression: Actual vs Predicted")
```


```
plot_results(y_test, y_pred_svr, "SVR: Actual vs Predicted")
```

**F)** The best model is SVR as it has the highest  $R^2$  score and the lowest MAE/MSE.

```
Multiple Linear Regression: MAE: 0.1281, MSE: 0.0276, R2: 0.7177  
Polynomial Regression: MAE: 0.1264, MSE: 0.1678, R2: -0.7161  
Support Vector Regression: MAE: 0.0997, MSE: 0.0207, R2: 0.7882
```



Polynomial Regression: Actual vs Predicted



SVR: Actual vs Predicted

